

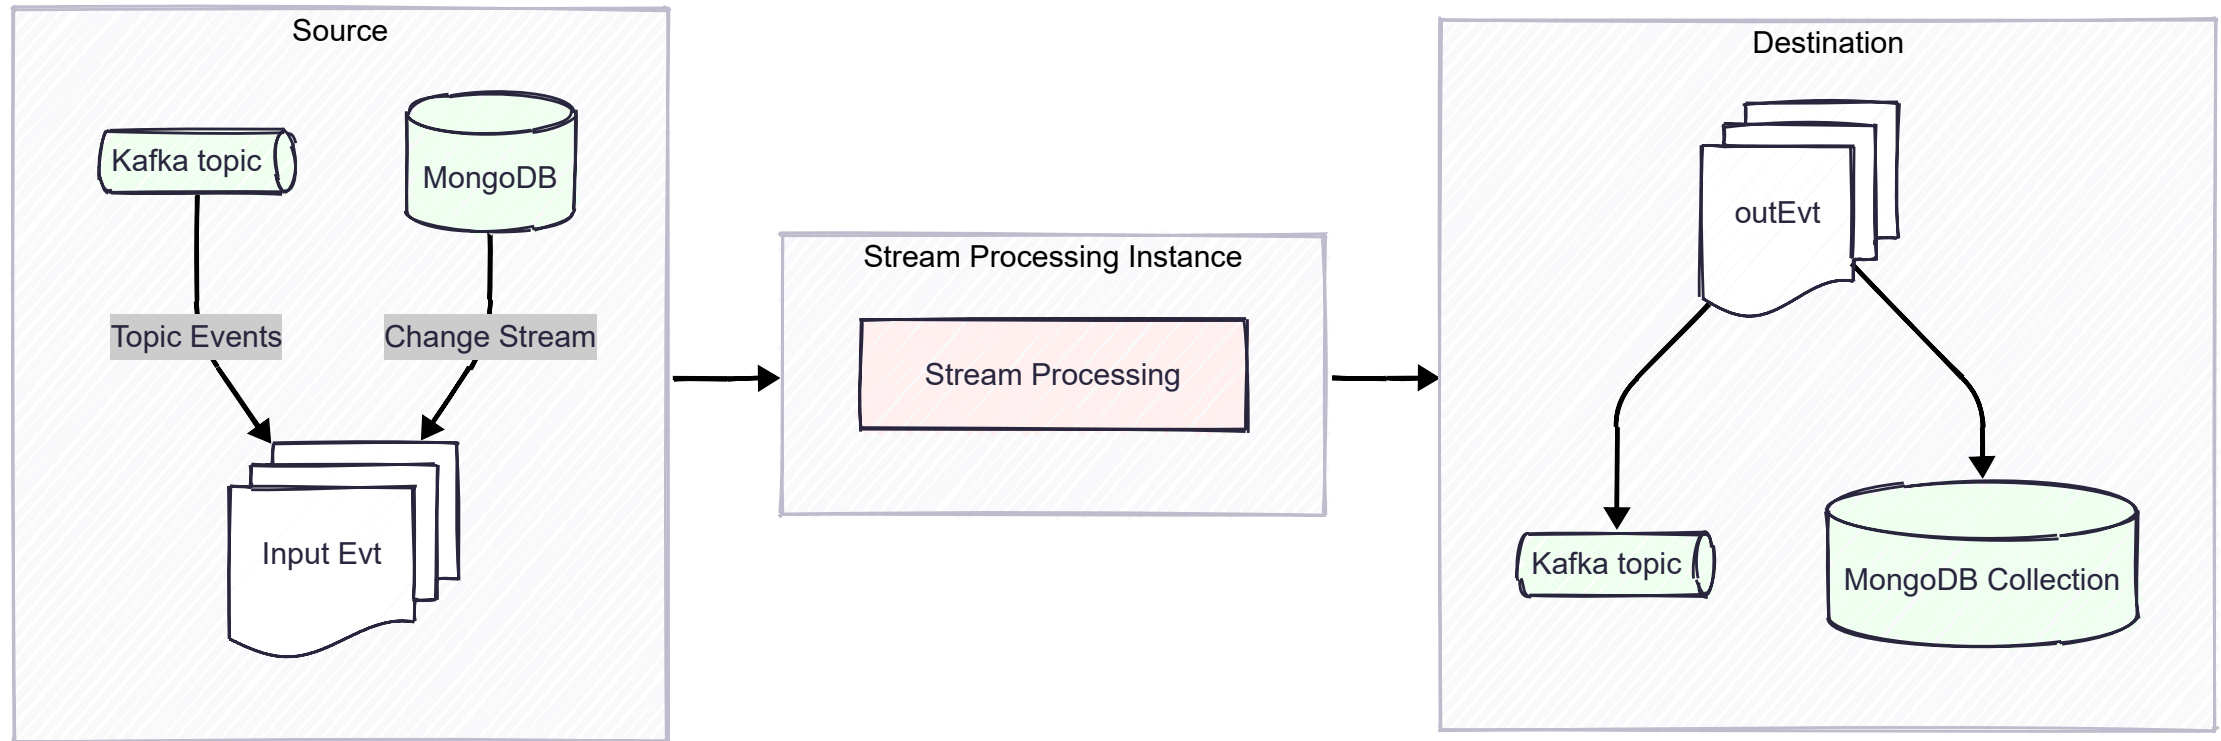
Events at the Movies with Atlas Stream Processing

A talk about streams of events, and
making sense of them.

Why?

- Event Driven Architecture
 - Lots of events, sporadic
 - Common integration pattern, existing
- Stream Analytics
 - Want to learn "what's going on"
 - "Recent Analytics", without the hotspots
 - No extra durable storage.

Processing as a Pipeline



Stream processor consumes events, and produces documents

In / Out

Connection Registry

- Kafka
- Atlas Database
- S3
- ...

Configure stream-processing-service

[Connect](#)[Stream Processors](#)[Monitoring](#)[Connection Registry](#)[+ Add connection](#)

Connection Name	Connection Type	Network Type 🔗	Actions
IoT_events	Atlas Database	Atlas Managed	✎ 🗑 🔗
click_buy_events	Atlas Database	Atlas Managed	✎ 🗑 🔗
kfk_1_topic_912	Apache Kafka	Atlas Managed	✎ 🗑 🔗
Dashboard_Receive_Q	Atlas Database	Atlas Managed	✎ 🗑 🔗

Stream Processor Connections

connectionName as configured

- 1st stage: \$source
- last state: \$merge | \$emit | \$externalFunction

```
[
  { $source: {
    connectionName: "mdbIn",
    db: "stream-demo",
    collection: "things" }},
  // { some processing stages...},
  { $merge:{
    into:{
      connectionName:"mdbConn",
      db:"db1",
      coll:"c1"} } }
]
```

Windowing and time basis

Windows are fixed width (usually).



Output is computed on events within its time boundaries.

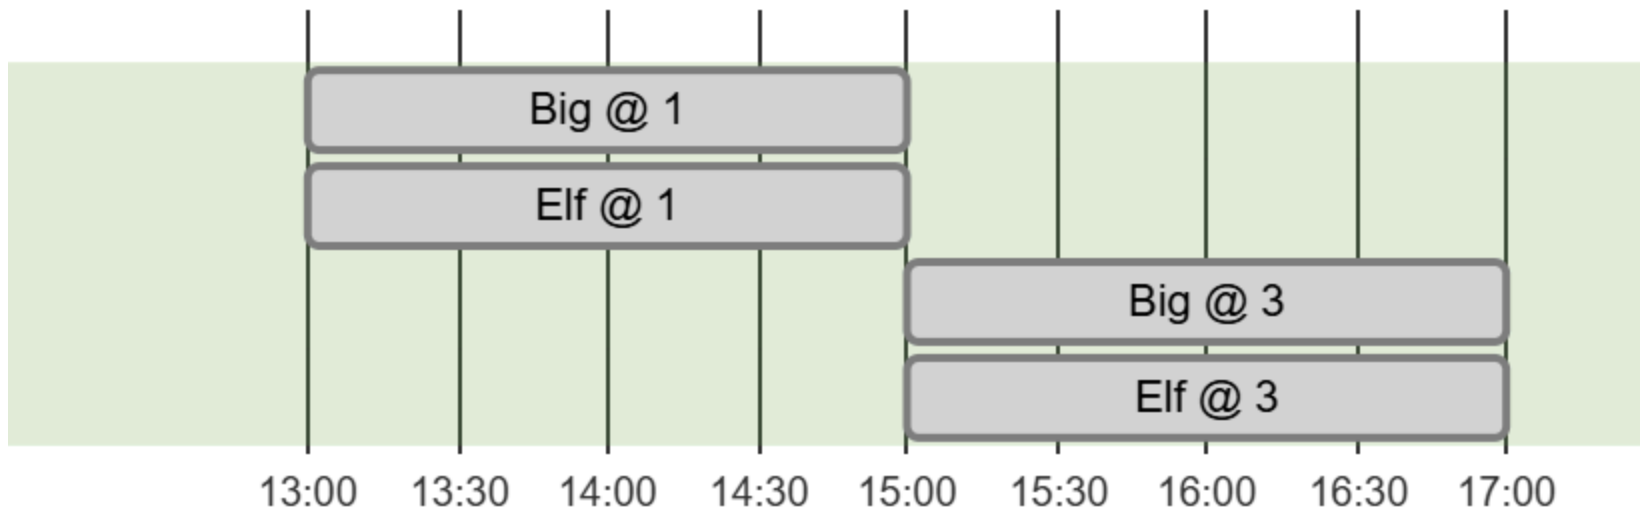
Create Stream Processor - How?

```
const pipeline = [{ $source: ... }, ...];  
sp.createStreamProcessor("mySP", pipeline)
```

- `pipeline` always starts with a `$source` stage.


Tumbling Window

Sequential Showtimes




Tumbling - How?

```
{
  $tumblingWindow: {
    interval: { size: 30, unit: "seconds" },
    pipeline: [
      {
        $group: {
          _id: "$movie",
          total: { $sum: "$amountPaid" }
        }
      }
    ] ...
  }
}
```




Hopping Window

Staggered Showtimes




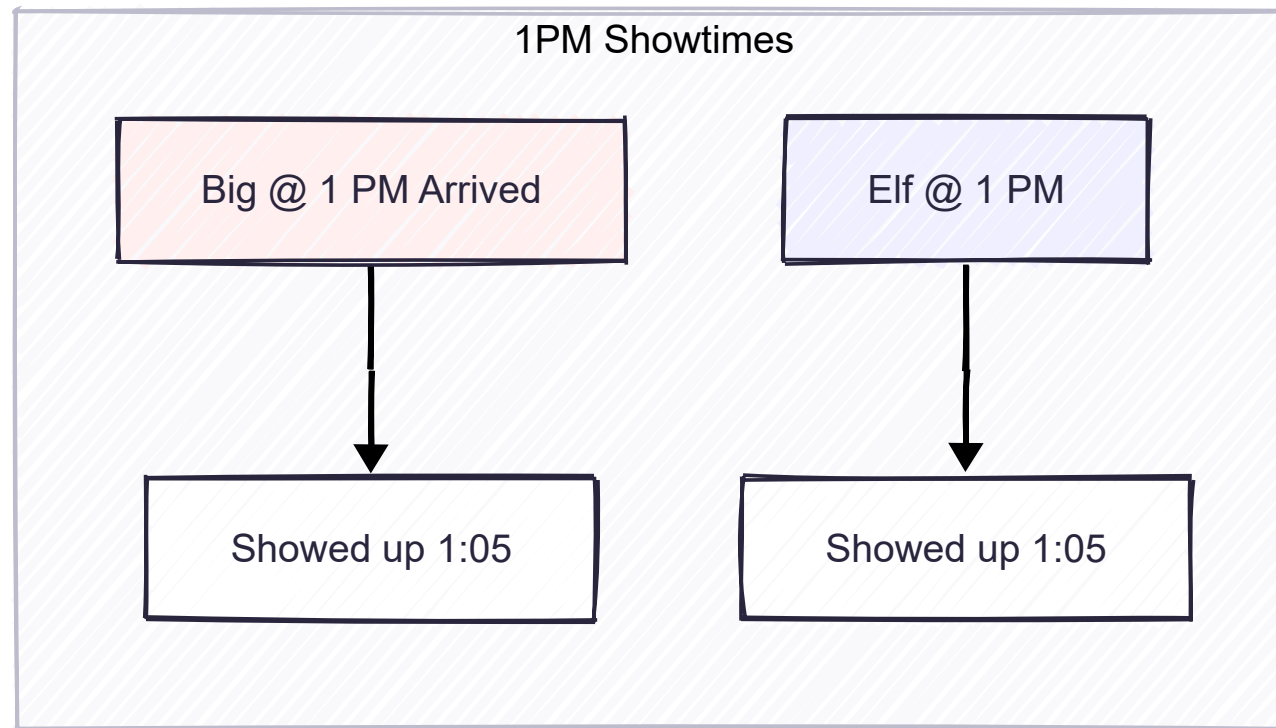
Hopping - How?

```
{
  $hoppingWindow:
  {
    interval: {size: 20, unit: "minute" },
    hopSize:  {size: 10, unit: "minute" },
    pipeline: [
      {
        $group: {
          _id: "$movie",
          walkIns: { $sum: "$ticketCount" }
        }
      }
    ] ...
  }
}
```




Missed the Window

Oops! What to do?

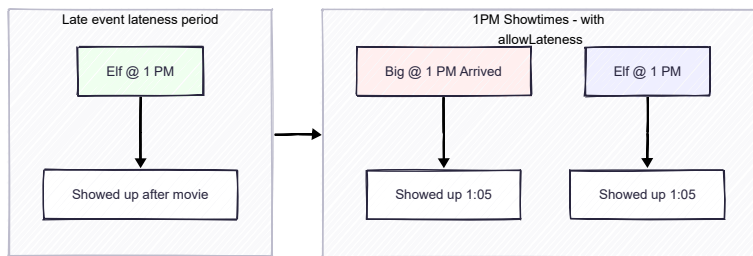


Allowed Lateness

`allowedLateness` lets late arrivals to be counted after *window-end-time*.




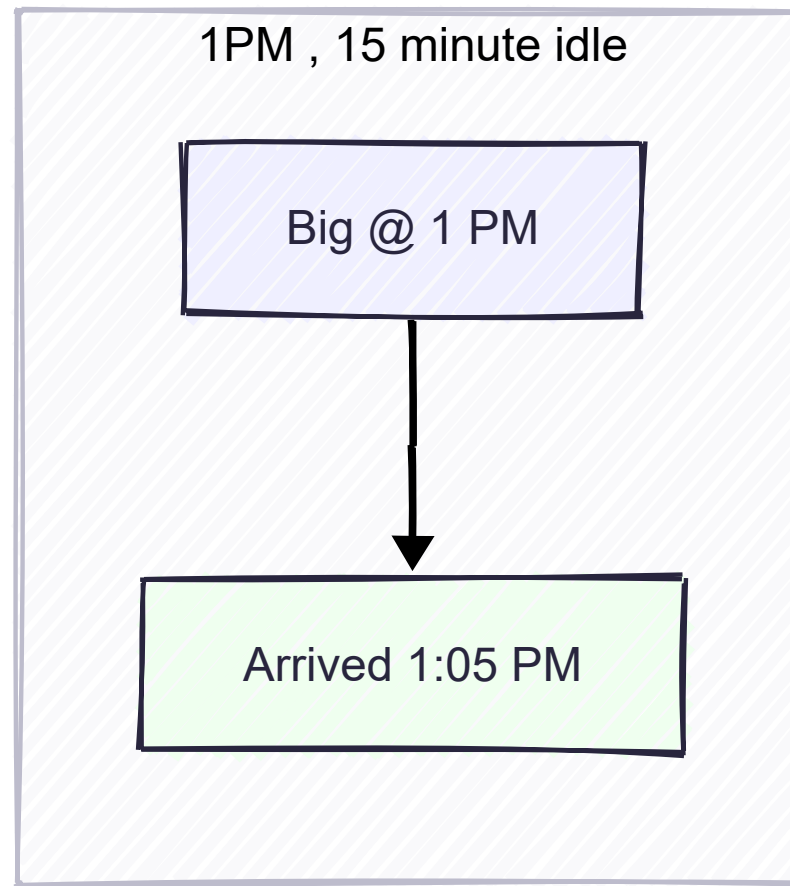
Lateness - How?



```
{
  $tumblingWindow: {
    allowedLateness: { size: 1, unit: "minute"},
    interval: { size: 30, unit: "seconds" },
    pipeline: [
      { $group: {
        _id: "$movie",
        total: { $sum: "$amountPaid" }
      } }
    ]
  }
}
```


Idle Time

Close the lobby
early - show in
progress



Late Event Handling

What happens when an event shows up **after** the window is closed?




Dead Letter Queue

What ends up in DLQ?

- Malformed
 - \$validate rejections
 - Payload deserialization errors
- Time Boundary Violations (late/early)
- Aggregation pipeline stage errors
- Full Document not available (change stream)

DLQ - How?

```
const options = {  
  dlq: {  
    connectionName: "my_dlq",  
    db: "my_db",  
    coll: "events_for_review"  
  }  
}  
  
sp.createStreamProcessor("mySP", /** pipeline */, options);
```



Session Window

- Closes when no event seen gap time after latest.

Thank You

MongoDB Champions

Nuri Halperin

LinkedIn: @nurih

nuri@plusnconsulting.com

